

1 Glint: VM image distribution in an OpenStack 2 multi-cloud environment

3 A. Charbonneau, M. Conklin, R. Demarais, C. Driemel, I. Gable,
4 C. Leavett-Brown, M. Paterson, R.J Sobie, R. Taylor

5 University of Victoria, Victoria, Canada

6 E-mail: rsobie@uvic.ca

7 **Abstract.** The use of cloud computing is becoming widespread in the HEP community. In
8 many instances, individual clouds are being federated to appear as a single infrastructure that
9 can be connected, for example, to the WLCG. For a small number of clouds, it is relatively
10 easy to manage the virtual machine images on all the sites. However, as the number of clouds
11 increases, then keeping track of the images requires a more robust management system. Glint
12 is designed to manage virtual machine images on OpenStack clouds. Glint makes it easy to
13 distribute images on multiple clouds in a reliable and error free manner.

14 1. Introduction

15 The use of clouds in HEP is becoming a significant source of computing resources for MC
16 production and data analysis. There are a number of compelling reasons for the migration to
17 clouds: clouds give researchers quick and dynamic access to unused or opportunistic resources;
18 sites find it easier to manage multiple project that have their own specific software requirements;
19 and cloud technology is supported by a large international community of developers.

20 If a cloud site is colocated with other traditional HEP computing resources or the resource
21 has been transformed to a cloud, then it can be viewed as another grid site on the WLCG.
22 Utilizing opportunistic resources (private non-HEP or commercial) can be done by directly
23 linking the cloud to the project workload management system (for example, VMDIRAC [1]).
24 Alternatively, one can unify the clouds into a single infrastructure (“grid of clouds”). The
25 HTCondor/CloudScheduler [2] or VAC/VCycle [3] were presented at this conference and are
26 two methods for unifying clouds.

27 We have established a distributed cloud system using HTCondor/CloudScheduler for particle
28 physics and astronomy applications [4, 5]. The distributed cloud is currently used by the ATLAS
29 [6] and Belle-II [7] experiments. The system is designed to boot user or project-specific virtual
30 machines (VM’s), making it easy to run HTC batch workloads from multiple projects.

31 The majority of the clouds today are using OpenStack software. OpenStack clouds require
32 that a user boot an image from the local Glance repository. If a user wants to use multiple
33 clouds, either manually or with our HTCondor/CloudScheduler system, then the image must
34 be manually uploaded to the Glance repository of each cloud. Transferring VM images manually
35 becomes tedious and error prone as one gets access to more IaaS clouds. We realized that we
36 required a better way of managing the VM images and we developed Glint help us manage VM
37 images over multiple OpenStack clouds.

The Glint service is designed for OpenStack clouds and the Glance repository with a pluggable architecture to allow support for different cloud types. Through a web browser or command-line interface, the user can identify clouds and control the distribution of images. Glint is currently used with the distributed cloud systems for ATLAS and Belle-II. In this paper, we describe the design and implementation of Glint. The source code is available from the Git repository [8].

2. Design

Glint is designed to operate as an OpenStack service. The Glint service is comprised of 4 components. The service component, which actually pushes and pulls data from different Glance repositories; The GUI component, which provides a user friendly interface integrated with OpenStack's Horizon dashboard; The Command Line Interface component, which allows users to try out glint without having to change or add another dashboard; And the Application Programming Interface, which allow developers to create their own tools. Figure ?? depicts a software component model of glint and it's integration with OpenStack.

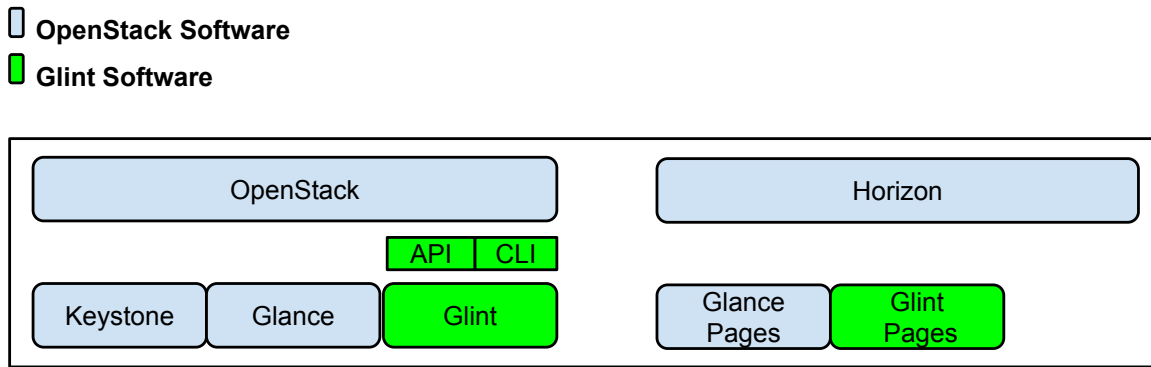


Figure 1. Glint Software Components

There are 3 key functions the Glint service provides; Glint allows users to create remote sites representing a remote glance repository; Glint users can add their remote repository credentials for Glint management; Finally, Glint manages image distribution as specified by the user. Image management is the most complicated operation provided by Glint. Figure 2 depicts the process of moving an image between two sites.

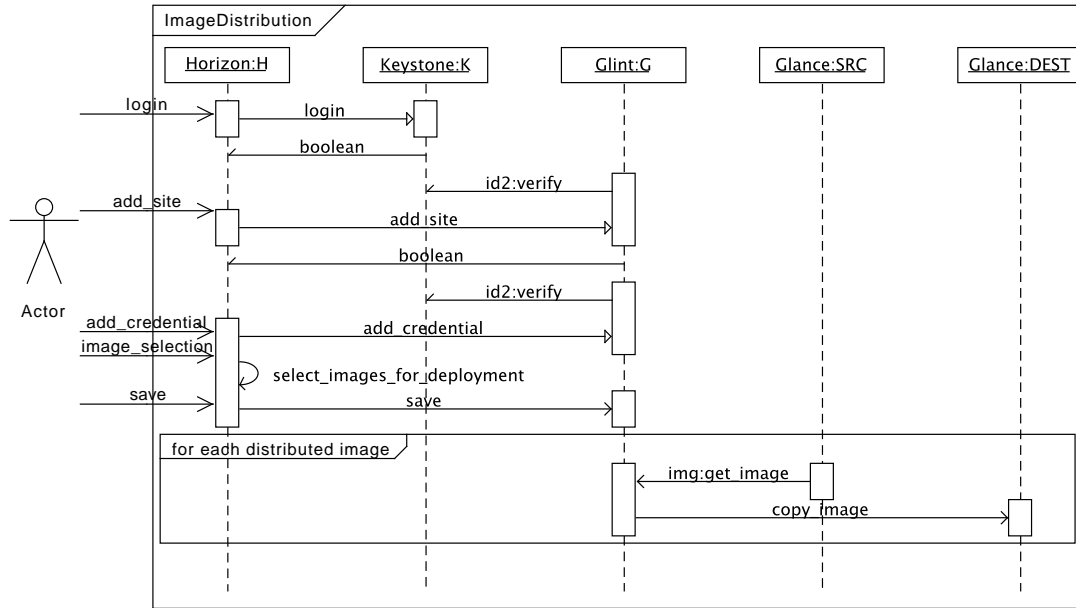


Figure 2. Image Distribution Design

56 Glint is designed around OpenStack's Keystone service for user authentication. This design
 57 requires a user to have an OpenStack account on the site that Glint is using as its primary host.
 58 Therefore, the user must login and register with Glint which will then use Keystone to check if
 59 the user is valid (c.f., Figure 3). If the user is valid, Keystone will generate an authentication
 60 token that the glint service persists. User requests are then performed on behalf of glint using
 61 the token for authentication.

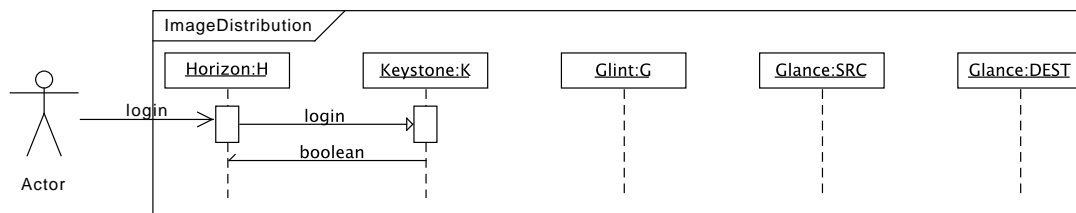


Figure 3. Glint User Login

62 The concepts of remote OpenStack sites and remote site credentials is essential in the design
 63 of glint. Users can add remote sites as a source and destination of images. Site credentials can
 64 then be added so Glint can use the site without having to request input from the user every
 65 time the site is accessed. Figure 4 depicts the user adding a remote site to the glint service.

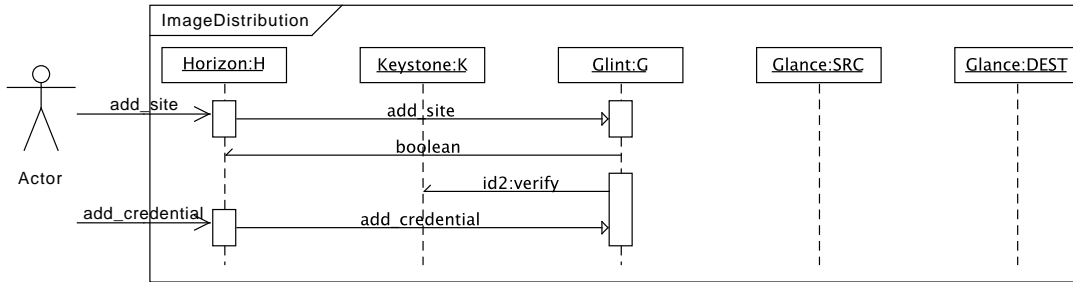


Figure 4. Glint - Create remote site and add remote credential

66 OpenStack's Glance service provides the image copying and deleting functionality. Glint relies
67 on this service to move images between remote sites. Figure 5 depicts the image distribution
68 process. It requires the user to select which images they want on each site, once the user is
69 happy they can save the new configuration. Glint proceeds to create a new thread for each
70 image transfer or delete image operation.

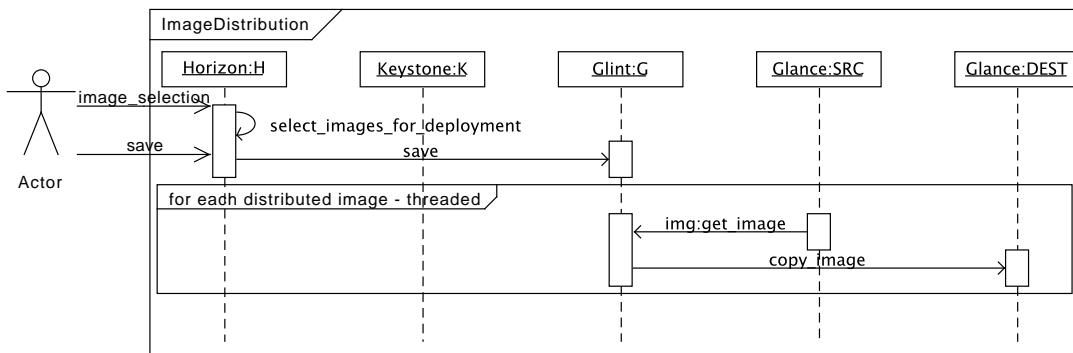


Figure 5. Glint - Reconfigure Image Distribution

3. Implementation

72 Glint design and implementation follow the OpenStack framework. There are 4 main components
73 developed and deployed to work with OpenStack. The image transfer component, the
74 Application Programming Interface (API), the Command Line Interface (CLI) and the horizon
75 web page elements.

76 The image transfer component is an HTTP service that uses Openstack's Keystone and
77 Glance services. The objective of the service is to integrate remote OpenStack clouds for the
78 purpose of moving images seamlessly among them. The implementation is written in the python
79 programming language and uses OpenStack's python APIs to integrate the Glance and Keystone
80 services.

81 The API is written in python and provides users of the Glint service to construct their own
82 applications using Glint. There are plans to provide Java and C/C++ APIs in the future.

83 The CLI is written in python and uses the API. The CLI is the first application to use the
84 Glint API. The CLI provides all the functionality that is provided by the modified Horizon
85 dashboard web pages added for Glint usage.

86 The modified Horizon dashboard is a set of added web pages to provide a friendly user
87 interface within the familiar Horizon dashboard. The pages are added to the images management
88 tab that Horizon already provides.

89 The code for the 4 main components is managed within 2 Eclipse projects using the PyDev
90 and Git plugins. The PyDev plugin understands the semantics of python and allows the Eclipse
91 Graphical User Interface (GUI) to provide proper functionality that is relevant for python
92 development. The Git plugin allows the use of Git and GitHub as our versioning system to
93 track changes and maintain the projects.

94 The first Eclipse project is a modified version of OpenStack's Horizon dashboard. This
95 dashboard includes pages relevant to image distribution. Web pages in Horizon are described
96 with python classes and use Django's powerful templating engine to provide a uniform
97 presentation to the user. There are two extra pages added to provide image distribution. The
98 first page allows users to add remote sites and add credentials to access remote sites. The second
99 page allows users to view all images on all sites, and provides an interface to add or remove
100 selected images from any site.

101 The second Eclipse project includes the API, CLI and the image transfer service. These are
102 developed in python and can use either Django or Paste to serve HTTP requests. The Django-
103 PyDev plugin is used for the project and provides Django related functions to manage the HTTP
104 service component of the software. An API is provided for python to extend the functionality
105 of the Glint service. The CLI is the first extension of the service that uses the API. The CLI is
106 written in python and takes advantage of the ArgParse python module which makes it easier to
107 develop CLI tools. The CLI and the modified Horizon interface are the two methods currently
108 offered to interact with the image distribution service. The CLI and the Horizon interface are
109 independent of each other, for example you can use the service with the CLI without having to
110 install the modified Horizon dashboard.

111 4. Deployment

112 4.1. Service Deployment

113 It is recommended to deploy the Glint service to the same secure machine that the Keystone and
114 Glance services are running. This is because glint maintains a database of user name, passwords
115 and authentication tokens in it's database. So it is important to keep glint on a secure machine.
116 Once glint is installed, you can use it immediately via the CLI.

117 Another deployment strategy is to also install the modified Horizon dashboard. This
118 dashboard allows users of the local OpenStack site to add remote sites specific for themselves
119 (i.e., other OpenStack sites the user has credentials on).

120 4.2. Image Deployment

121 Figure 6 depicts a basic deployment of an image using Glint with 3 OpenStack sites. Each
122 OpenStack site is independent and has their own user and tenant management. The credentials
123 of each site is stored on Glint securely and used by the service when the image deployment needs
124 to be changed. For example, the user requested a deployment change such that Image A is to
125 be duplicated from Site 1 to Site2 and Site3. Glint will first authenticate against Site1 using
126 Keystone, then perform a copy operation for Site 2 and Site 3. For a successful copy operation,
127 Glint will first authenticate the Site with credentials found in it's database for that site.

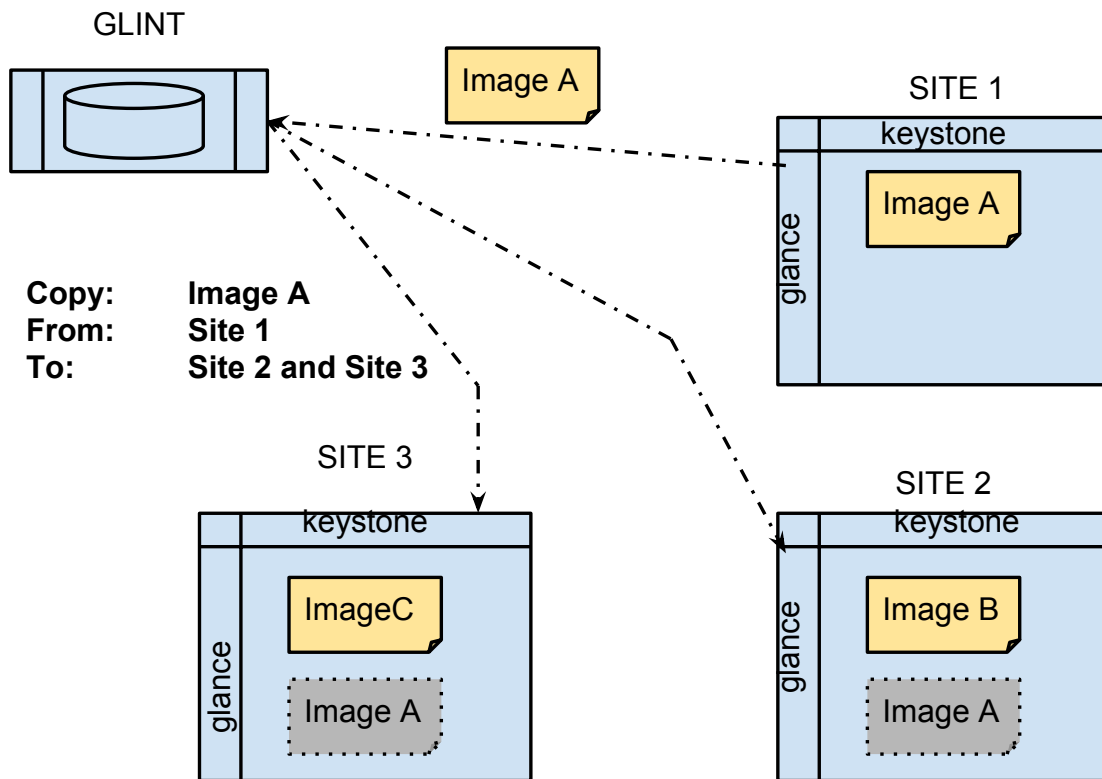


Figure 6. Image Distribution

Glint is available on PyDev as required for deployment via StackForge for OpenStack integration. The latest version is currently deployed for image distribution at the University of Victoria.

References

- [1] VM dirac
- [2] Gable chep talk
- [3] VAC vcycle
- [4] HPCS cloud paper
- [5] Sobie NYC talk
- [6] Ryan CHEP talk
- [7] Sobie Belle-II talk
- [8] *Glint*, a system for managing VM images on multiple clouds. <https://github.com/hep-gc/glint>